

A DATA-EFFICIENT UNIFIED MULTIMODAL FUSION FRAMEWORK FOR ART PRICE PREDICTION: AN AUCTION MARKET STUDY

1. INTRODUCTION

The art market is a complex domain where the valuation of artworks depends on factors such as visual aesthetics, artist reputation, historical significance, and market trends. Traditional methods rely on expert opinions and subjective judgment, which may lead to inconsistencies and a lack of transparency. To address these challenges, this project proposes a data-driven approach using machine learning and deep learning techniques for predicting artwork prices.

The system develops a unified multimodal fusion framework that combines visual and metadata features to predict artwork prices in auction markets. The dataset includes artwork images along with metadata such as artist name, year of creation, medium, and previous auction prices. Data preprocessing is performed using Python libraries like Pandas, including cleaning, normalization, and handling missing values.

The model is implemented using the PyTorch framework, where visual features are extracted using Convolutional Neural Networks (CNNs) with the support of TorchVision and PIL, while metadata features are encoded into structured formats. These features are fused into a unified model, and regression techniques are applied to predict artwork prices. The model performance is evaluated using metrics such as Mean Absolute Error (MAE), Root Mean Squared Error (RMSE), and the R-squared (R^2) score.

The system is backend-oriented and operates through Python scripts without a graphical user interface, with results displayed in the console. The proposed framework assists collectors, auction houses, and investors in making informed decisions. Overall, the project demonstrates that integrating multimodal data improves the accuracy and reliability of art price prediction.

1.1 Overview of the System

The proposed system is a data-driven art price prediction model that utilizes a multimodal fusion approach to estimate the prices of artworks in auction markets. The system integrates both visual features extracted from artwork images and metadata features such as artist name, year of creation, medium, and previous auction prices. It is implemented using Python and the PyTorch framework, with supporting libraries like Pandas for data preprocessing, TorchVision and PIL for image processing, and Scikit-learn for evaluation.

The system performs key operations, including data loading, preprocessing, feature extraction, model training, and price prediction. Visual features are extracted using Convolutional Neural Networks (CNNs), while metadata is structured and encoded for analysis. These features are combined in a unified model to improve prediction accuracy. The system is backend-oriented and operates through Python scripts, where the predicted results and evaluation metrics are displayed in the console without a graphical user interface. This work aims to improve prediction accuracy through a unified multimodal learning approach.

1.2 Objectives of the System

- To develop a data-driven approach for predicting artwork prices and reducing dependency on subjective and manual valuation methods.
- To integrate visual and metadata features using a multimodal fusion technique.
- To extract meaningful visual features using deep learning models (CNNs) and preprocess metadata for effective analysis.
- To build and train a regression model for accurate price prediction.
- To evaluate the model performance using metrics such as MAE, RMSE, and R^2 score.

1.3 Scope of the Project

The project's scope lies in providing a reliable and efficient system for predicting artwork prices using machine learning and deep learning techniques. The system can assist collectors, auction houses, and investors in making informed decisions based on data-driven insights. It focuses on combining multiple data modalities to improve prediction accuracy compared to traditional methods.

The project can be further extended by incorporating real-time auction data, advanced deep learning architectures, and a graphical user interface for better user interaction. Additionally, the system can be adapted to other domains where price prediction is required by modifying the dataset and feature inputs.

1.4 Dataset Description Table

The dataset used in this project contains artwork information along with structured metadata and image paths. The features used for model training are described below.

Feature	Description
price	Selling price of the artwork (target variable)
artist	Name of the artist
title	Title of the artwork
yearCreation	Year in which the artwork was created
signed	Information about whether the artwork is signed
condition	Physical condition of the artwork
period	Historical period (e.g., Contemporary, Post-War)
movement	Art movement (e.g., Abstract, Surrealism)
image_path	Path to the artwork image

2. SYSTEM ANALYSIS

2.1 Existing System

Traditional art price estimation systems primarily rely on expert opinions, historical auction records, and manual evaluation methods. These approaches consider factors such as artist reputation, artwork condition, and past sales data. In some cases, basic statistical or machine learning models are used, but they mainly depend on structured metadata and do not incorporate visual features of the artwork. As a result, these systems lack a comprehensive understanding of the artwork's aesthetic value and market dynamics.

2.2 Limitations of the Existing System

- Highly dependent on subjective human judgment.
- Lack of consistency and transparency in price estimation.
- Limited use of data-driven and automated techniques.
- Inability to utilize the visual features of artworks.
- Lower prediction accuracy due to reliance only on metadata.
- Not scalable for large and diverse datasets.

2.3 Proposed System

The proposed system introduces a data-driven multimodal fusion framework for predicting artwork prices. It integrates both visual features extracted from artwork images and metadata such as artist name, year, medium, and previous auction prices.

The system is implemented using Python and the PyTorch framework. Visual features are extracted using Convolutional Neural Networks (CNNs) with the help of Torchvision and PIL, while metadata is preprocessed using Pandas. These features are combined into a unified model, and regression techniques are applied to predict artwork prices. The system operates as a backend model and provides outputs through the console using deep learning-based multimodal fusion techniques.

2.4 Advantages of the Proposed System

- Provides a data-driven and objective approach to price prediction.
- Improves accuracy by combining visual and metadata features.
- Utilizes deep learning techniques for better feature extraction.
- Reduces human bias and manual errors.
- Scalable and adaptable to large datasets.
- Can be extended with advanced models and real-time data.

3. SYSTEM SPECIFICATIONS

3.1 Hardware Requirements

Component	Specification
Processor	Intel Core i5 or higher
RAM	Minimum 8 GB (16 GB recommended)
Storage	256 GB HDD/SSD
GPU (Optional)	NVIDIA GPU
System Type	64-bit system
Input Devices	Keyboard, Mouse
Output Device	Monitor

3.2 Software Requirements

The proposed system is implemented using a combination of software tools and libraries that support efficient multimodal data processing and deep learning. The development is carried out on platforms including Windows, Linux, and macOS, ensuring cross-platform compatibility.

Python is used as the primary programming language due to its extensive support for scientific computing and machine learning applications. The deep learning framework employed is PyTorch, which provides dynamic computation graphs and efficient model training capabilities.

Pandas is utilized for structured data handling and preprocessing of metadata, while NumPy supports high-performance numerical computations. Scikit-learn is used for preprocessing and evaluation, enabling performance assessment through standard metrics.

For image processing tasks, Torchvision is employed for data transformations and feature extraction, and the Python Imaging Library (PIL) is used for image loading and manipulation.

The implementation and experimentation are conducted using development environments such as Jupyter Notebook, Visual Studio Code, and PyCharm, which facilitate efficient coding, debugging, and visualization.

4. SYSTEM ARCHITECTURE

4.1 Module Description

The proposed system is a **backend-driven application** that predicts artwork prices using a **multimodal deep learning model**, integrating both visual and metadata features for enhanced accuracy. It follows a modular architecture comprising data preparation, feature extraction, multimodal fusion, and price prediction.

The **data preparation module** ensures clean, structured inputs, while the **feature extraction module** captures meaningful representations from images (via CNNs) and metadata. These features are combined in the **multimodal fusion stage**, enabling the model to learn complex relationships across data types. The **price prediction module** then applies regression techniques to generate precise estimations.

All modules operate cohesively to deliver efficient processing, robust feature learning, and reliable predictions. The system's modular and scalable design makes it adaptable for real-world deployment in artwork valuation.

4.2 Data Preparation Module

The Data Preparation Module is responsible for collecting, organizing, and preprocessing the dataset used for training and testing the model. The dataset consists of artwork images along with metadata attributes such as artist name, year of creation, medium, and previous auction prices.

This module performs several preprocessing tasks to improve data quality. These include handling missing values, removing inconsistencies, and normalizing the data to bring it into a standard format. Metadata features are cleaned and converted into structured numerical representations suitable for machine learning models. Image data is also prepared by resizing, formatting, and converting into tensors using libraries such as PIL and Torchvision.

The preprocessing steps are implemented using Python libraries like Pandas and NumPy, ensuring that the dataset is clean, consistent, and ready for further processing.

4.3 Feature Extraction Module

The Feature Extraction Module focuses on extracting meaningful and representative features from both image and metadata inputs. For visual data, Convolutional Neural Networks (CNNs) are used to automatically learn important patterns, such as textures, shapes, and artistic styles, from artwork images.

The system utilizes PyTorch and TorchVision to implement CNN-based feature extraction, while PIL is used for loading and handling images. These visual features capture the aesthetic aspects of the artwork, which play a crucial role in price determination.

For metadata, categorical and numerical features such as artist name, year, and medium are encoded into structured formats. These features provide contextual information that complements the visual data. The extracted features from both sources are then prepared for integration in the next stage.

4.4 Multimodal Fusion Module

The Multimodal Fusion Module integrates the features obtained from both visual and metadata sources into a unified representation. This module plays a key role in improving the model's prediction capability by combining different types of information.

The fusion process ensures that both visual characteristics (from images) and contextual details (from metadata) are considered simultaneously during model training. This integration allows the system to capture complex relationships between different features that cannot be achieved using a single data source.

By combining multiple modalities, the model gains a more comprehensive understanding of the factors influencing artwork prices, leading to improved accuracy and robustness.

4.5 Price Prediction Module


The Price Prediction Module is the final stage of the system, where the fused features are used to estimate the price of artworks. A regression-based model is trained using the combined feature set to predict continuous price values.

The model is implemented using PyTorch and is trained on the prepared dataset. During training, the model learns the relationship between input features and artwork prices. Once trained, it can predict prices for new or unseen data.

The performance of the model is evaluated using standard evaluation metrics such as Mean Absolute Error (MAE), Root Mean Squared Error (RMSE), and R-squared (R^2) scores. These metrics help in assessing the accuracy and reliability of the model.

The system operates through Python scripts, and the predicted results, along with evaluation metrics, are displayed in the console, as the project does not include a graphical user interface.

4.6 System Architecture Diagrams



5. IMPLEMENTATION

5.1 Implementation Procedure

The implementation of the proposed system follows a structured approach using Python and the PyTorch framework. Initially, the dataset consisting of artwork images and metadata is collected and organized. The data is then preprocessed using libraries such as Pandas, where missing values are handled, and the data is normalized for consistency.

Next, image data is processed using PIL and TorchVision, where images are resized and converted into tensors suitable for deep learning models. Metadata features are encoded into structured numerical formats. A Convolutional Neural Network (CNN) model is implemented using PyTorch to extract visual features from images.

The extracted visual features and encoded metadata features are then combined using a multimodal fusion technique. The fused data is passed into a regression model, which is trained to predict artwork prices. Finally, the model is evaluated using performance metrics such as MAE, RMSE, and R^2 score.

5.2 Execution of Modules

The system executes each module sequentially to ensure proper data flow and processing.

- The **Data Preparation Module** loads the dataset and performs preprocessing tasks such as cleaning, normalization, and handling missing values.
- The **Feature Extraction Module** processes image data using a CNN and encodes metadata into structured formats.
- The **Multimodal Fusion Module** integrates both visual and metadata features into a unified representation.
- The **Price Prediction Module** uses the fused features to train the regression model and predict artwork prices.

The entire execution is carried out through Python scripts, and outputs are generated in the console.

6. RESULTS AND DISCUSSION

6.1 Model Performance

The performance of the model is evaluated using standard regression metrics to ensure accuracy and reliability.

- **Mean Absolute Error (MAE):** Measures the average difference between predicted and actual prices.
- **Root Mean Squared Error (RMSE):** Evaluates the magnitude of prediction errors.
- **R-squared (R²) Score:** Indicates how well the model explains the variance in the data.

Model Evaluation Results

Model Type	MAE	RMSE	R ² Score
Image Only	7498.54	14744.39	-1.1676
Metadata Only	6318.38	11841.25	-0.3981
Text Only	6319.62	11841.91	-0.3982
Simple Fusion	7253.19	18025.73	-2.2398
Attention Fusion	5661.92	10010.97	0.0007

The model performance is evaluated across different approaches, such as image-only, metadata-only, text-only, simple fusion, and attention-based fusion models, to compare their effectiveness.

6.2 Result Analysis

The experimental results indicate that single-modality models, including image-only, metadata-only, and text-only approaches, exhibit higher error values, leading to lower prediction accuracy. This is because each modality captures only partial information about the artwork, limiting the model's overall understanding.

The multimodal fusion approach improves prediction performance by integrating both visual and metadata features, thereby providing a more comprehensive representation of the data. However, simple fusion techniques are not sufficiently effective in capturing complex interdependencies between features.

The attention-based fusion model outperforms all other approaches by dynamically assigning importance to different features. This results in lower error values and improved prediction accuracy, demonstrating the effectiveness of attention mechanisms in learning meaningful cross-modal relationships.

Overall, the results confirm that integrating multiple data modalities significantly enhances the accuracy and reliability of artwork price prediction.

7. CONCLUSION AND FUTURE ENHANCEMENT

7.1 Conclusion

The project successfully implements a data-driven system for predicting artwork prices using a multimodal deep learning approach. By integrating visual and metadata features, the system overcomes the limitations of traditional subjective valuation methods. The use of PyTorch and CNN enables effective feature extraction, while regression techniques ensure accurate predictions. The system is efficient, scalable, and provides reliable results, making it useful for real-world applications in the art market. The proposed system provides a scalable and reliable solution for real-world art market applications.

Furthermore, the system highlights the effectiveness of combining multiple data modalities to enhance prediction performance. It reduces dependency on manual expertise and promotes a more objective and consistent valuation process. The project also demonstrates the practical application of machine learning and deep learning techniques in solving real-world problems. Overall, it lays a strong foundation for future advancements in AI-based art valuation systems.

7.2 Future Enhancement

- Integration of real-time auction and market trend data.
- Implementation of advanced deep learning models for improved accuracy.
- Development of a graphical user interface for better user interaction.
- Expansion of the dataset to include more diverse artworks.
- Deployment of the system as a web or cloud-based application.
- Optimization of model performance using hyperparameter tuning.

8. BIBLIOGRAPHY

- [1] B. Li and T. Liu, “An Analysis of Multi-Modal Deep Learning for Art Price Appraisal,” in *Proc. IEEE Int. Conf. on Parallel and Distributed Processing with Applications, Big Data and Cloud Computing, Sustainable Computing and Communications, Social Computing and Networking (ISPA/BDCloud/SocialCom/SustainCom)*, 2021, pp. 1509–1513, doi: 10.1109/ISPA-BDCloud-SocialCom-SustainCom52081.2021.00203.
- [2] C. Liu, “Prediction and Analysis of Artwork Price Based on Deep Neural Network,” *Scientific Programming*, vol. 2022, Article ID 7133910, 2022, doi: 10.1155/2022/7133910.
- [3] P. Mauer and S. Paszkiel, “Tabular Data Models for Predicting Art Auction Results,” *Applied Sciences*, vol. 14, 11006, 2024, doi: 10.3390/app142311006.
- [4] K. Lee, J. Park, S. Goree, D. Crandall, and Y.-Y. Ahn, “Social Signals Predict Contemporary Art Prices Better than Visual Features, Particularly in Emerging Markets,” *Scientific Reports*, vol. 14, 11615, 2024, doi: 10.1038/s41598-024-60957-z.
- [5] J. Park, Y. Lee, D. Yang, J. Park, and H. Jung, “Artwork Pricing Model Integrating the Popularity and Ability of Artists,” *Advances in Statistical Analysis*, vol. 108, pp. 889–913, 2024, doi: 10.1007/s10182-024-00504-3.

9. APPENDIX

9.1 Sample Code Snippet:

```
import pandas as pd

from PIL import Image

from torchvision import transforms


# Load dataset

data = pd.read_csv("dataset.csv")


# Define image transformations

transform = transforms.Compose([

    transforms.Resize((224, 224)),

    transforms.ToTensor()])


# Load and preprocess an image

img_path = data['image_path'][0]

image = Image.open(img_path).convert("RGB")

image = transform(image)

print(image.shape) # Output: [3, 224, 224]
```

9.2 Model Evaluation Snippet:

```
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

import numpy as np

# Example predictions

y_true = np.array([10000, 15000, 20000])

y_pred = np.array([9500, 14000, 21000])

mae = mean_absolute_error(y_true, y_pred)

rmse = np.sqrt(mean_squared_error(y_true, y_pred))

r2 = r2_score(y_true, y_pred)

print("MAE:", mae)

print("RMSE:", rmse)


print("R2 Score:", r2)
```

9.3 Output Sample:

The system generates predicted artwork prices using multiple modeling approaches, including simple multimodal fusion and attention-based fusion techniques. The performance is evaluated using k-fold cross-validation and metrics such as Mean Absolute Error (MAE), Root Mean Squared Error (RMSE), and R-squared (R^2) scores.

The results show that the attention-based fusion model outperforms the simple fusion approach by effectively capturing cross-modal relationships and assigning importance to different features.

The output of the simple fusion model is shown in **Fig. A.1**, where features are combined using basic concatenation.



```
(venv) PS D:\FINAL_DATASET> python src/train_model_simple.py
Epoch [18/20] Train Loss: 4.9366
Epoch [19/20] Train Loss: 4.9826
Epoch [20/20] Train Loss: 4.3849
Fold 4 Results -> MAE: 7290.08, RMSE: 14027.16, R2: -0.7528

===== Fold 5 =====
D:\FINAL_DATASET\venv\lib\site-packages\torchvision\models\_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed
in the future, please use 'weights' instead.
  warnings.warn(
D:\FINAL_DATASET\venv\lib\site-packages\torchvision\models\_utils.py:223: UserWarning: Arguments other than a weight enum or 'None' for 'weights' are deprecate
d since 0.13 and may be removed in the future. The current behavior is equivalent to passing 'weights=EfficientNet_B0_Weights.DEFAULT'. You can also us
e 'weights=EfficientNet_B0_Weights.DEFAULT' to get the most up-to-date weights.
  warnings.warn(msg)
Epoch [1/20] Train Loss: 58.8280
Epoch [2/20] Train Loss: 42.2251
Epoch [3/20] Train Loss: 27.0029
Epoch [4/20] Train Loss: 16.4235
Epoch [5/20] Train Loss: 11.2011
Epoch [6/20] Train Loss: 9.6133
Epoch [7/20] Train Loss: 8.2829
Epoch [8/20] Train Loss: 7.7854
Epoch [9/20] Train Loss: 7.5752
Epoch [10/20] Train Loss: 7.5144
Epoch [11/20] Train Loss: 6.9325
Epoch [12/20] Train Loss: 6.7810
Epoch [13/20] Train Loss: 6.4977
Epoch [14/20] Train Loss: 5.7612
Epoch [15/20] Train Loss: 5.6592
Epoch [16/20] Train Loss: 5.5093
Epoch [17/20] Train Loss: 4.8706
Epoch [18/20] Train Loss: 4.3989
Epoch [19/20] Train Loss: 4.4981
Epoch [20/20] Train Loss: 4.3970
Fold 5 Results -> MAE: 11979.83, RMSE: 41670.25, R2: -3.0249

===== Final Average Performance =====
Average MAE: 10436.47
Average RMSE: 41237.16
Average R2: -23.0013
(venv) PS D:\FINAL_DATASET>
```

Fig. A.1 Simple Fusion Output

The performance of the model evaluated using k-fold cross-validation is illustrated in **Fig. A.2**, ensuring the robustness of the model.

```
(venv) PS D:\FINAL_DATASET> python src\train_model_KFold.py

Epoch [11/20] Train Loss: 3.9497
Epoch [12/20] Train Loss: 3.6826
Epoch [13/20] Train Loss: 3.2807
Epoch [14/20] Train Loss: 3.2848
Epoch [15/20] Train Loss: 3.1195
Epoch [16/20] Train Loss: 2.5339
Epoch [17/20] Train Loss: 2.4626
Epoch [18/20] Train Loss: 2.5396
Epoch [19/20] Train Loss: 2.2748
Epoch [20/20] Train Loss: 1.9990
Fold 4 Results + MAE: 5919.57, RMSE: 11563.73, R2: -0.1912


===== Fold 5 =====
Epoch [1/20] Train Loss: 62.3201
Epoch [2/20] Train Loss: 54.8722
Epoch [3/20] Train Loss: 38.4790
Epoch [4/20] Train Loss: 18.4170
Epoch [5/20] Train Loss: 11.2956
Epoch [6/20] Train Loss: 9.9025
Epoch [7/20] Train Loss: 9.2583
Epoch [8/20] Train Loss: 7.5779
Epoch [9/20] Train Loss: 6.9281
Epoch [10/20] Train Loss: 5.4014
Epoch [11/20] Train Loss: 4.6521
Epoch [12/20] Train Loss: 4.1234
Epoch [13/20] Train Loss: 3.8102
Epoch [14/20] Train Loss: 3.5299
Epoch [15/20] Train Loss: 3.2852
Epoch [16/20] Train Loss: 2.8935
Epoch [17/20] Train Loss: 2.6150
Epoch [18/20] Train Loss: 2.8958
Epoch [19/20] Train Loss: 2.4139
Epoch [20/20] Train Loss: 2.3132
fold 5 Results + MAE: 7676.21, RMSE: 21612.77, R2: -0.0827

===== Final Average Performance =====
Average MAE: 5721.76
Average RMSE: 12648.59
Average R2: -0.0877

(venv) PS D:\FINAL_DATASET>
```

Fig. A.2 K-Fold Validation Output

The results of the proposed attention-based fusion model are presented in **Fig. A.3**, showing improved prediction performance.



The screenshot shows a Visual Studio Code interface with a terminal window open. The terminal displays the output of a Python script named `src\evaluate.py`. The output shows the results of model evaluation for different fusion methods. The attention-based fusion model shows the best performance with the lowest MAE and RMSE and the highest R2.

```
PS D:\FINAL_DATASET> & d:\FINAL_DATASET\venv\Scripts\Activate.ps1
(venv) PS D:\FINAL_DATASET> python src\evaluate.py

===== MODEL EVALUATION RESULTS =====

Image Only
MAE : 7498.54
RMSE : 14744.39
R2 : -1.1676

Metadata Only
MAE : 6318.38
RMSE : 11841.25
R2 : -0.3981

Text Only
MAE : 6319.62
RMSE : 11841.91
R2 : -0.3982

Simple Fusion
MAE : 7253.19
RMSE : 18025.73
R2 : -2.2398

Attention Fusion
MAE : 5661.92
RMSE : 10010.97
R2 : 0.0007

(venv) PS D:\FINAL_DATASET>
```

Fig. A.3 Attention-Based Model Output

A sample of predicted artwork prices compared with true values and predictions from different models is shown in **Fig. A.4**. It can be observed that attention-based predictions are closer to the true price.

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R
1	True Price	Image Prediction	Metadata Prediction	Text Prediction	Simple Fusion Prediction	Attention Fusion Prediction												
2	1275	3100.8794	0.027860375	0.045018546	1094.3292	518.08594												
3	680.00006	215.5408	0.46498194	0.059956536	221.1109	560.65204												
4	9499.997	6086.2524	1.1536732	0.045018546	2261.985	2975.9934												
5	4999.9995	556.9803	1.2789595	0.080270074	4600.174	1890.0005												
6	800.0001	453.89502	0.5525439	0.045018546	539.9309	1322.2299												
7	15000.004	2863.6084	0.6552781	0.054329704	4321.6997	2859.9023												
8	1275	14.635212	0.6117121	0.1443012	112.663795	117.75034												
9	800.0001	10286.842	0.12598923	0.1443012	6047.8213	10360.433												
10	9350.004	498.7052	0.80464745	0.10851838	1050.9067	1555.8162												
11	1499.9998	842.58435	0.12060964	0.045018546	349.99356	1110.0958												
12	1275	96.05662	1.0332458	0.1443012	211.34584	288.82968												
13	900.0001	260.77988	0.44484055	0.063101985	370.14084	1062.4464												
14	800.0001	985.1527	0.27002513	0.044999786	2864.3408	2846.715												
15	680.00006	856.44196	0.62712085	0.045018546	1131.6338	1442.0837												
16	800.0001	3723.2935	0.41312262	0.045018546	3866.1497	2676.2915												
17	4500.0005	149.79073	0.5080024	0.09341566	413.43198	321.00702												
18	8000.0005	2.5368319	0.1297391	0.09906169	12.337225	48.546917												
19	4999.9995	2628.8972	-0.040883146	0.045018546	2837.4153	6800.4824												
20	7499.999	0.21488732	0.7410456	0.05550508	13.795747	24.847126												
21	5999.9995	540.1657	0.9429679	0.06740646	791.1134	1257.2708												
22	4999.9995	44995.51	0.08960111	0.09490797	18966.404	10602.6875												
23	7000.001	129.93657	0.56233144	0.045018546	214.4278	317.44284												
24	15000.004	37.68895	0.3678669	0.070619084	127.13301	218.81296												
25	4999.9995	785.6525	0.32162392	0.045018546	1953.3318	1168.9824												
26	800.0001	0.99247503	0.58453035	0.1443012	1.8914106	6.1521263												
27	800.0001	1449.7883	0.12645747	0.092099145	2307.051	5729.8647												
28	1275	936.81537	0.92851496	0.089376956	1036.5132	921.2058												
29	141.60646	0.3368561	0.045018546	0.045018546	415.6380	678.04136												

Fig. A.4 Predicted Output

A sample of the dataset used in the project is shown in **Fig. A.5**, including artwork prices along with structured metadata and corresponding image paths used for multimodal learning.

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P
	Unnamed: 0	price	artist	title	year	Creation	signed	condition	period	movement	image_path					
1	0	28,500 USD	Tommaso Otti	Bayreuth Opera	2021	Signed on vers	This work is in e	Contemporary	Baroque		images/auction_images/0.jpg					
2	1	3,000 USD	Pavel Tchelit	Drawings of the O	First Half 20th C	Signed and titl	Not examined	c	Post-War	Surrealism	images/auction_images/1.jpg					
3	2	5,000 USD	Leo Gabin	Two on Sidewalk	2016	Signed, titled	c	This work is in e	Contemporary	Abstract	images/auction_images/2.jpg					
4	3	5,000 USD	Matthias Dorn	Blumenszene	2010	Signed, titled	c	This work is in e	Contemporary	Abstract	images/auction_images/3.jpg					
5	4	2,500 USD	Alexis Margue	Feverish Embarkat	2001	Signed on vers	c	This work is in e	Contemporary	Abstract	images/auction_images/4.jpg					
6	5	7,575 USD	Kenzo Okada	Bamboo	1977	Signed lower r	This work is in e	Contemporary	Abstract Expressionism		images/auction_images/5.jpg					
7	6	7,550 USD	Francesco Clei	Air	2007	Numbered an	This work is in e	Contemporary	Neo-Expressionism		images/auction_images/6.jpg					
8	7	3,550 USD	GÅxñther FÅ	Untitled (Green)	1993	Signed lower r	This work is in v	Contemporary	Modernism		images/auction_images/7.jpg					
9	8	3,075 USD	Dan Walsh	Manifold - Blue	2014	Signed lower r	This work is in e	Contemporary	Minimalism		images/auction_images/8.jpg					
10	9	3,550 USD	GÅxñther FÅ	Untitled (Orange/E	1993	Signed lower r	This work is in e	Contemporary	Modernism		images/auction_images/9.jpg					
11	10	3,075 USD	Dan Walsh	Manifold - Red	2014	Signed lower r	This work is in e	Contemporary	Minimalism		images/auction_images/10.jpg					
12	11	40,000 USD	Friedel Dzabas	Upstream	1973	Signed and da	This work is in e	Contemporary	Abstract		images/auction_images/11.jpg					
13	12	15,000 USD	Zhang Yu	96.6.2	[nan]	Signed lower r	This work is in v	Contemporary	[nan]		images/auction_images/12.jpg					
14	13	7,000 USD	Chris Ofili	Last Night, New D	2008	Signed lower r	This work is in e	Contemporary	Punk, Young British Artists, Abstract		images/auction_images/13.jpg					
15	14	4,000 USD	Marcel Dzama	Here's a Fine Revo	2015	Signed lower r	This work is in e	Contemporary	Magic Realism		images/auction_images/14.jpg					
16	15	4,000 USD	Leonardo Drex	CPP11	2015	Signed lower r	This work is in e	Contemporary	Post-Minimalism		images/auction_images/15.jpg					
17	16	5,000 USD	Anne Appleby	Cottonwood	2012	Signed lower r	This work is in e	Contemporary	Minimalism		images/auction_images/16.jpg					
18	17	8,000 USD	Sol LeWitt	Horizontal Bands (	2002	Signed lower r	This work is in e	Contemporary	Minimalism		images/auction_images/17.jpg					
19	18	5,500 USD	Jockum Nordst	House and Bugs	2008	Signed lower r	This work is in e	Contemporary	Surrealism		images/auction_images/18.jpg					
20	19	2,500 USD	William Bailey	Stradina	2002	Signed lower r	This work is in e	Contemporary	Realism		images/auction_images/19.jpg					
21	20	7,000 USD	Al Held	Fly Away	1992	Signed verso	This work is in e	Contemporary	Abstract Expressionism		images/auction_images/20.jpg					
22	21	6,000 USD	Tommaso Abts	Untitled (small circ	2015	Signed on low	The work is in e	Contemporary	Abstract		images/auction_images/21.jpg					
23	22	7,000 USD	Richard Tuttle	Trans Asian	1993	Signed lower r	This work is in e	Contemporary	Post-Minimalism		images/auction_images/22.jpg					
24	23	7,000 USD	Matt Mullican	REPRESENT THE W	2020	Signed lower r	This work is in e	Contemporary	Conceptual		images/auction_images/23.jpg					
25	24	8,000 USD	Leonardo Drex	CPP1	2015	Signed lower r	This work is in e	Contemporary	Post-Minimalism		images/auction_images/24.jpg					
26	25	6,500 USD	Pat Steir	Hand Painted Lily #1992 - 2004	Signed lower r	This work is in e	Contemporary	Conceptual			images/auction_images/25.jpg					
27	26	10,000 USD	Wayne Thieba	Dark Canyon Bluff	2014	Signed bottom	This work is in e	Contemporary	Modernism		images/auction_images/26.jpg					
28	27	40,000 USD	Arish Kone	Untitled 12	1991	Signed lower r	This work is in e	Contemporary	Conceptual		images/auction_images/27.jpg					

Fig. A.5 Dataset Sample